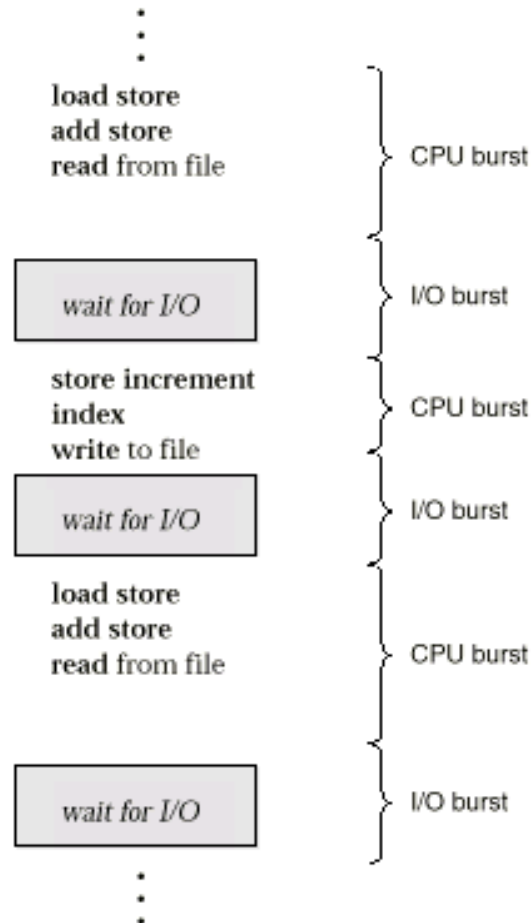


Operating Systems

CPU Scheduling

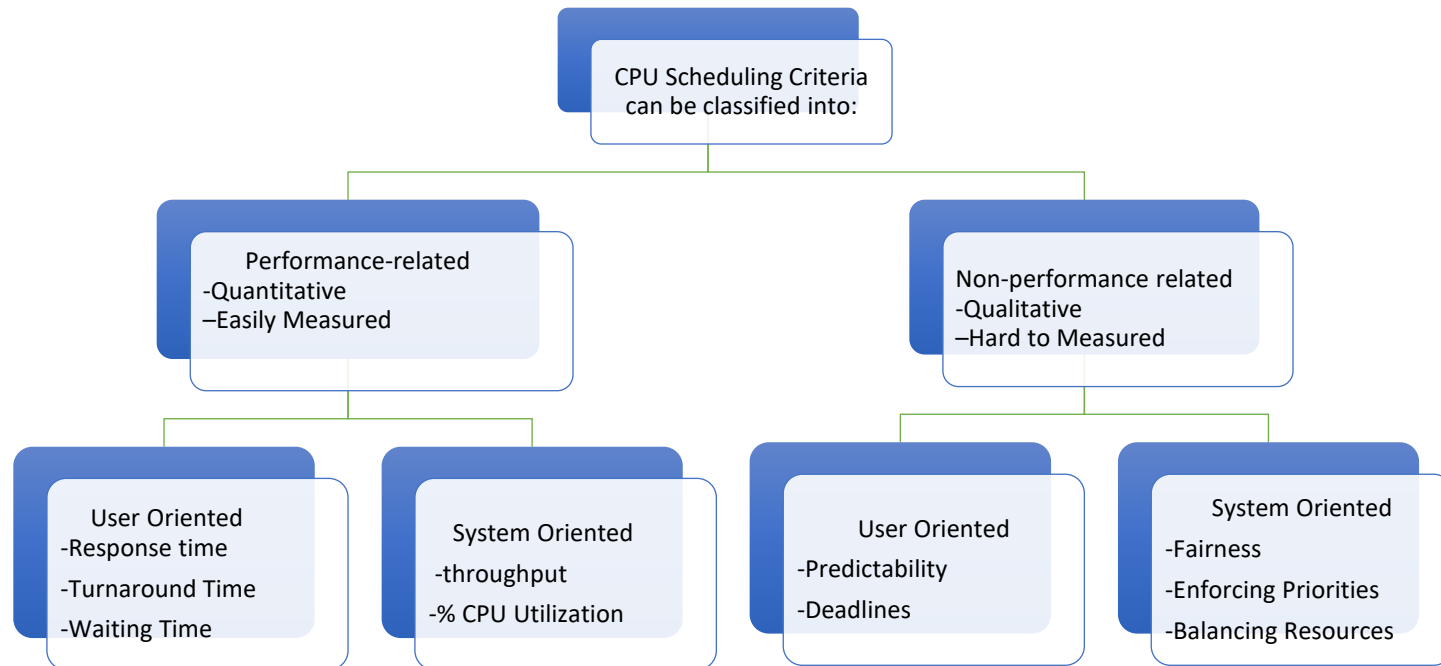
- Fundamentally, scheduling is a matter of managing queues to minimize queuing delay and to optimize performance in a queuing environment.
- Scheduling needs to meet system objectives, such as:
 - ✓ minimize response time
 - ✓ maximize throughput
 - ✓ maximize processor efficiency
 - ✓ support multiprogramming
- Scheduling is central to OS design



“CPU-bound”
processes require
more CPU time than
I/O time

“I/O-bound”
processes spend
most of their time
waiting for I/O.

- **Uniprocessor Scheduling:** scheduling a *single* CPU among all the processes in the system



- Turnaround time
 - ✓ Time from submission to completion (batch Process)
 - ✓ CPU-bound process equivalent of response time
- Response time
 - ✓ Time to start responding (interactive users)
 - ✓ I/O bound processes
- Throughput
 - ✓ Number of Process processed per unit of time
 - ✓ Productivity
- Processor(CPU) utilization
 - ✓ Percent of time CPU is busy
 - ✓ Efficiency
- Waiting time
 - ✓ Time spent waiting in READY Queue
 - ✓ Each process should get a fair share of the CPU

- Predictability
 - ✓ Same time/cost regardless of load on the system
- Deadlines
 - ✓ Maximize number of deadlines met
- Fairness
 - ✓ No process should suffer starvation
- Enforcing priorities
 - ✓ Favor higher priority processes
- Balancing resources
 - ✓ Keep system resources busy

- All systems
 - ✓ Fairness: give each process a fair share of the CPU
 - ✓ Enforcement: ensure that the stated policy is carried out
 - ✓ Balance: keep all parts of the system busy
- Batch systems
 - ✓ Throughput: maximize jobs per unit time (hour)
 - ✓ Turnaround time: minimize time users wait for jobs
 - ✓ CPU utilization: keep the CPU as busy as possible
- Interactive systems
 - ✓ Response time: respond quickly to users' requests
 - ✓ Proportionality: meet users' expectations
- Real-time systems
 - ✓ Meet deadlines: missing deadlines is a system failure!
 - ✓ Predictability: same type of behavior for each time slice

- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them
 - ✓ When does scheduler make decisions? When a process :
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is **non-preemptive**
 - ✓ Processes keep CPU until it releases either by terminating or I/O wait.
- All other scheduling is **preemptive**
 - ✓ Interrupts

- Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ✓ switching context
 - ✓ switching to user mode
 - ✓ Pumping to the proper location in the user program to restart that program
- **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running

- Important in scheduling and resource allocation algorithms
- Policy
 - ✓ What is to be done
- Mechanism
 - ✓ How to do it
- Policy: All users equal access
- Mechanism: round robin scheduling
- Policy: Paid Ps get higher priority
- Mechanism: Preemptive scheduling algorithm

- Dispatcher
 - ✓ Low-level: mechanism
 - ✓ Responsibility: context switch
- Scheduler
 - ✓ High-level: policy
 - ✓ Responsibility: deciding which process to run
- Could have an allocator for CPU as well
 - ✓ Parallel and distributed systems

Selection function

- Which process in the ready queue is selected next for execution?

Decision mode

- At what times is the selection function exercised?
 - ✓ **Nonpreemptive**
 - ❑ A process in the running state runs until it blocks or ends
 - ✓ **Preemptive**
 - ❑ Currently running process may be interrupted and moved to the Ready state by the OS
 - ❑ Prevents any one process from monopolizing the CPU

Process	Arrival Time	Burst Time
1	0	3
2	2	6
3	4	4
4	6	5
5	8	2


- **Selection function:** the process that has been waiting the longest in the ready queue (hence, FCFS, FIFO queue)
- **Decision mode:** non-preemptive- a process runs until it blocks itself (I/O or other)
- Pros:
 - ✓ Simple to implement
 - ✓ Fair policy
- Cons:
 - ✓ Favors CPU-bound processes
 - ❑ A process that does not perform any I/O will monopolize the processor!
 - ❑ I/O-bound processes have to wait until CPU-bound process completes
 - ❑ They may have to wait even when their I/Os have completed
 - poor device utilization
 - ❑ We could reduce the average wait time by giving more priority to I/O bound processes
 - ✓ waiting time depends on arrival order
 - ✓ **Convoy effect** - short process stuck waiting for long process

P #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	3	9	13	18	20
TAT/Response time	3	7	9	12	12
Service time	3	6	4	5	2
Wait time	0	1	5	7	10

Average response time = $(3+7+9+12+12)/5 = 8.6$

Throughput = $5/20 = 1/4$

Ave (RespTime/ServTime) = $(1+7/6+9/4+12/5+12/2)/5 = 2.564$




- **Selection function:** the process with the shortest expected CPU burst time
- **Decision mode:** non-preemptive
- SJF implicitly incorporates priorities: shortest Jobs are given preference.
 - ✓ Typically these are I/O bound Ps
- Lack of preemption not suitable in a time sharing environment
 - ✓ CPU bound process gets lower priority
 - ✓ But a process doing no I/O at all could monopolize the CPU if it is the first one in the system
- **Advantages:**
 - ✓ Minimizes average wait time. Provably optimal
- **Disadvantages**
 - ✓ Not practical: difficult to predict burst time
 - ❑ Possible: past predicts future
 - ✓ May starve long jobs

P #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	3	9	15	20	11
TAT/Response time	3	7	11	14	3
Service time	3	6	4	5	2
Wait time	0	1	7	9	1

Average response time = $(3+7+11+14+3)/5 = 7.6$

Throughput = $5/20 = 1/4$

Ave (RespTime/ServTime) = $(1+7/6+11/4+14/5+3/2)/5 = 1.844$



- Can average all past history equally
- But recent history of a process is more likely to reflect future behavior
- A common technique for that is to use *exponential averaging*

1. t_n = actual length of n^{th} CPU burst
2. τ_{n+1} = predicted value for the next CPU burst
3. $\alpha, 0 \leq \alpha \leq 1$
4. Define :

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n.$$

- ✓ Puts more weight on recent instances whenever $\alpha > 1/n$

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:

$$\tau_{n+1} = \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots$$

$$+ (1 - \alpha)^p \alpha t_{n-p} + \dots$$

$$+ (1 - \alpha)^{n+1} \tau_0$$
- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor


- Shortest Remaining Time (SRT) = Preemptive SJF
- If a process arrives in the Ready queue with estimated CPU burst less than remaining time of the currently running process, **preempt**.
- Prevents long Process's from dominating.
 - ✓ But must keep track of remaining burst times
- Better turnaround time than SJF
 - ✓ Short Process's get **immediate** preference

P #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	3	15	8	20	10
TAT/Response time	3	13	4	14	2
Service time	3	6	4	5	2
Wait time	0	7	0	9	0

Average response time = $(3+13+4+14+2)/5 = 7.2$

Throughput = $5/20 = 1/4$

Ave (RespTime/ServTime) = $(1+13/6+1+14/5+1)/5 = 1.6$




- **Selection function:** same as FCFS
- **Decision mode:** Preemptive
- Each process gets a small unit of CPU time (*time quantum*), usually 10-100 milliseconds.
 - ✓ After this time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are n processes in the ready queue and the time quantum is q , then each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
 - ✓ No process waits more than $(n-1)q$ time units.

P #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	4	18	17	20	15
TAT/Response time	4	16	13	14	7
Service time	3	6	4	5	2
Wait time	1	10	9	9	5

Average response time = $(4+16+13+14+7)/5 = 10.8$

P P S S S S S S S S S S S S S S Switch to P4

Round
Robin
T = 1




P #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	5	17	13	20	15
TAT/Response time	5	15	9	14	7
Service time	3	6	4	5	2
Wait time	2	9	5	9	5


P2 e P P P P P P2 t P4 t P3 f P5 f P2 fi P4 time expires
 Que Que Q Q Q Q Q Que Que Que Que Que Queue: P4
 P1 s Swit S S P S P Swit Swit Swit Swit Swit P4 Continues

Average response time = $(5+15+9+14+7)/5 = 10$

Round
 Robin
 T = 2



- Must be substantially larger than process switch time
- Should be larger than the typical CPU burst
- If too large, degenerates to FCFS
- Too small, excessive context switches (overhead)



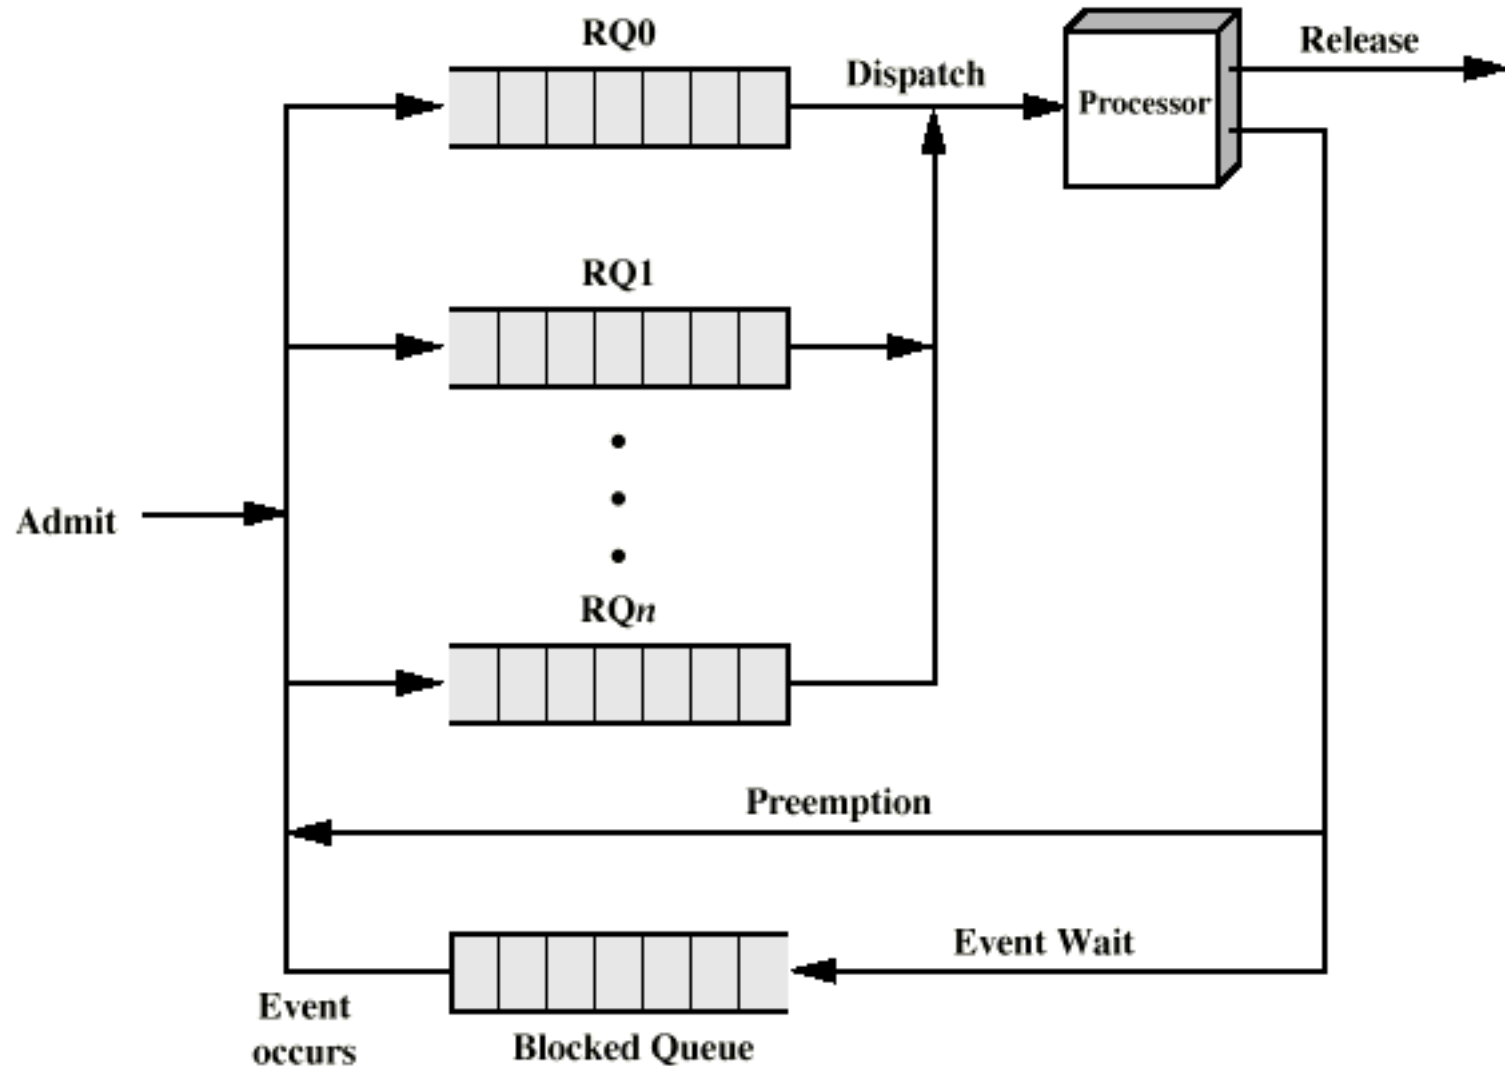
- Each context switch has the OS using the CPU instead of the user process
 - ✓ give up CPU, save all info, reload w/ status of incoming process
 - ✓ Say 20 ms quantum length, 5 ms context switch
 - ✓ Waste of resources
 - ❑ 20% of CPU time (5/20) for context switch
 - ✓ If 500 ms quantum, better use of resources
 - ❑ 1% of CPU time (5/500) for context switch
 - ❑ Bad if lots of users in system – interactive users waiting for CPU
 - ✓ Balance found depends on P mix

- Advantages
 - ✓ Low response time, good interactivity
 - ✓ Fair allocation of CPU across processes
 - ✓ Low average waiting time when job lengths vary widely
- Disadvantages
 - ✓ Poor average waiting time when jobs have similar lengths
 - ❑ Average waiting time is even worse than FCFS!
 - ✓ Performance depends on length of time slice
 - ✓ Still favours CPU-bound processes
 - ❑ An I/O bound process uses the CPU for a time less than the time quantum and then is blocked waiting for I/O
 - ❑ A CPU-bound process runs for its whole time slice and goes back into the ready queue (in front of the blocked processes)

- One solution: virtual round robin (VRR)
 - ✓ When a I/O has completed, the blocked process is moved to an auxiliary queue which gets *preference* over the main ready queue
 - ✓ A process dispatched from the auxiliary queue gets a shorter time quantum (what is “left over” from its quantum when it was last selected from the ready queue)

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority)
 - ✓ Preemptive
 - ✓ nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem \equiv Starvation – low priority processes may never execute
- Solution \equiv Aging – as time progresses increase the priority of the process

Scheduler chooses from low priority queue only if higher ones



- Non-preemptive, tries to get best average normalized turnaround time
- Choose next process with the highest ratio

$$\text{Ratio} = \frac{\text{time spent waiting} + \text{expected service time}}{\text{expected service time}}$$

- Attractive approach to scheduling because it accounts for the age of a process.
- While shorter processes are favored (a smaller denominator yields a larger ratio), aging without service increases the ratio so that a longer process will eventually get past competing shorter processes .


$$RR = (\text{waiting-time} + \text{service-time}) / \text{service-time}$$

Process #	1	2	3	4	5
Arrival time	0	2	4	6	8
Finish time	3	9	13	20	15
TAT/Response time	3	7	9	14	7
Service time	3	6	4	5	2
Wait time	0	0	7	9	0

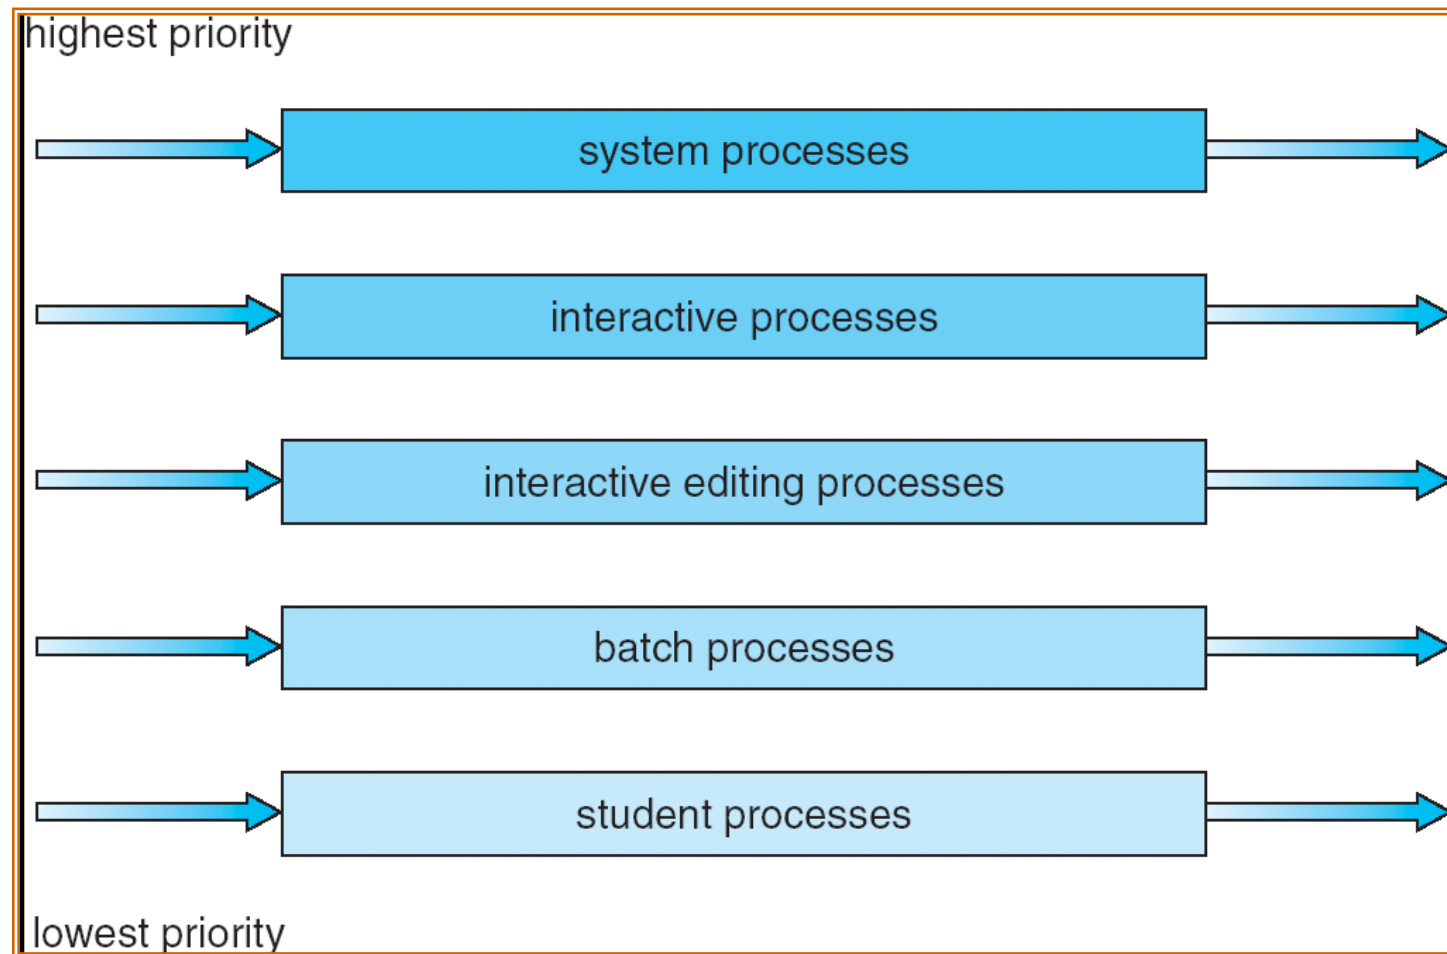
Average response time = $(3+7+9+14+7)/5 = 8$

Throughput = $5/20 = 1/4$


Ave (RespTime/ServTime) = $(1+7/6+9/4+14/5+7/2)/5 = 2.144$



- Ready queue is partitioned into separate queues:
foreground (interactive)
background (batch)
- Each queue has its own scheduling algorithm
 - ✓ foreground – RR
 - ✓ background – FCFS
- Scheduling must be done between the queues
 - ✓ Fixed priority scheduling; (i.e., serve all from foreground then from background). Possibility of starvation.
 - ✓ Time slice – each queue gets a certain amount of CPU time which it can schedule amongst its processes; i.e., 80% to foreground in RR
 - ✓ 20% to background in FCFS




- Several READY queues with decreasing priorities:
 - ✓ $P(RQ0) > P(RQ1) > \dots > P(RQn)$
 - ✓ But lower priority queues get longer time slice
- New processes are placed in RQ0
- If they use their time quantum, they are placed in RQ1. If they time out again, they go to RQ2, etc. until they reach lowest priority
- Automatically reduces priority of CPU-bound Ps, leaving I/O-bound ones at top priority
- Dispatcher always chooses a process from highest non-empty queue
- Problem: long Ps can “starve”
 - ✓ Solution: “aging” (promote priority of a process that waits too long in a lower priority queue)



- Note: many variations of policies possible
✓ including different policies at different levels..

- Three queues:
 - ✓ Q_0 – RR with time quantum 8 milliseconds
 - ✓ Q_1 – RR time quantum 16 milliseconds
 - ✓ Q_2 – FCFS
- Scheduling
 - ✓ A new P enters queue Q_0 which is served FCFS. When it gains CPU, P receives 8 milliseconds. If it does not finish in 8 milliseconds, P is moved to queue Q_1 .
 - ✓ At Q_1 P is again served FCFS and receives 16 additional milliseconds. If it still does not complete, it is preempted and moved to queue Q_2 .



- Multilevel feedback scheduling is the most flexible method
- Can customize:
 - ✓ Number of queues
 - ✓ Scheduling algorithm per queue
 - ✓ Priority “upgrade” criteria
 - ✓ Priority “demotion” criteria
 - ✓ Policy to determine which queue a process starts in

Bibliography

- ❖ Silberschatz, A, Galvin, P.B, and Gagne, G., Operating System Principles, 9e, John Wiley & Sons, 2013.
- ❖ Stallings W., Operating Systems-Internals and Design Principles, 7e, Pearson Education, 2014.
- ❖ Harvey M. Deital, “Operating System”, Third Edition, Pearson Education, 2013.
- ❖ Andrew S. Tanenbaum, “Modern Operating Systems”, Second Edition, Pearson Education, 2004.
- ❖ Gary Nutt, “Operating Systems”, Third Edition, Pearson Education, 2004.

Acknowledgements

- ❖ I have drawn materials from various sources such as mentioned in bibliography or freely available on Internet to prepare this presentation.
- ❖ I sincerely acknowledge all sources, their contributions and extend my courtesy to use their contribution and knowledge for educational purpose.

Thank You!!

?